

Class08

Dahlia Loomis

4/23/23

1. Exploratory Data Analysis

First, we will read the data.

```
setwd("/Users/dahlialoomis/Desktop/WisconsinCancer")
# Save your input data file into your Project directory
fna.data <- "WisconsinCancer.csv"
# Complete the following code to input the data and store as wisc.df
wisc.df <- read.csv(fna.data, row.names = 1)
```

Now, I am examining the data to make sure that column names are set correctly.

```
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
--	---------------	------------------------	-----------	------------	--------------

842302	0.2419		0.07871	1.0950	0.9053	8.589
842517	0.1812		0.05667	0.5435	0.7339	3.398
84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.1560	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
area_se smoothness_se compactness_se concavity_se concave.points_se						
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
84358402	94.44	0.011490	0.02461	0.05688		0.01885
843786	27.19	0.007510	0.03345	0.03672		0.01137
symmetry_se fractal_dimension_se radius_worst texture_worst						
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
84358402	0.01756		0.005115	22.54		16.67
843786	0.02165		0.005082	15.47		23.75
perimeter_worst area_worst smoothness_worst compactness_worst						
842302	184.60	2019.0		0.1622		0.6656
842517	158.80	1956.0		0.1238		0.1866
84300903	152.50	1709.0		0.1444		0.4245
84348301	98.87	567.7		0.2098		0.8663
84358402	152.20	1575.0		0.1374		0.2050
843786	103.40	741.6		0.1791		0.5249
concavity_worst concave.points_worst symmetry_worst						
842302	0.7119		0.2654		0.4601	
842517	0.2416		0.1860		0.2750	
84300903	0.4504		0.2430		0.3613	
84348301	0.6869		0.2575		0.6638	
84358402	0.4000		0.1625		0.2364	
843786	0.5355		0.1741		0.3985	
fractal_dimension_worst						
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

```
#looks good. The ID is the row name
#diagnosis is the first column
```

Now, we are removing the first diagnosis column so that it is not present in the data set.

```
#Use -1 to remove the first column
wisc.data <- wisc.df[,-1]
```

Set up a new vector called diagnosis

```
diagnosis <- wisc.df[,1]
diag <- as.factor(diagnosis)
```

Let's explore the data set:

- **Q1.** How many observations are in this data set?

```
#we can use the nrow()
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations.

- **Q2.** How many of the observations have a malignant diagnosis?

```
#We can use the table() command
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

There are 212 observations that have a malignant diagnosis.

- Q3.** How many variables/features in the data are suffixed with `_mean`?

```
.mean <- grep("_mean", colnames(wisc.data))

length(.mean)
```

[1] 10

There are 10 variables that are suffixed with `_mean`.

2. Principal Component Analysis (PCA)

First we will check to see if the data need to be scaled before we perform PCA.

```
# Check column means and standard deviations
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean

7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Since the columns are in different units, this indicates that scaling is necessary.

Now, we will apply PCA.

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

```
wisc.pr <- prcomp(wisc.data, scale = TRUE)

summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005

Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

From the summary function, the proportion of the original variance captured by PC1 was 0.4427.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

We need three PCs to describe at least 70 percent of the data.

```
pca.var <- wisc.pr$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per[1]
```

```
[1] 44.3
```

```
pca.var.per[1]
```

```
[1] 44.3
```

```
pca.var.per[1] + pca.var.per[2]
```

```
[1] 63.3
```

```
pca.var.per[1] + pca.var.per[2] + pca.var.per[3]
```

```
[1] 72.7
```

```
#sum = 0
#for (i in 1:length(pca.var.per)){
#  add = pca.var.per[i]
#  sum = sum + add
#  if (sum > 0.7) {
```

```
# print(i)
# }
#}
```


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs are required to describe at least 90% of the original variance in the data.


Interpreting PCA Results

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

```
plot(wisc.pr) #generates a barplot, which is not what I want
```

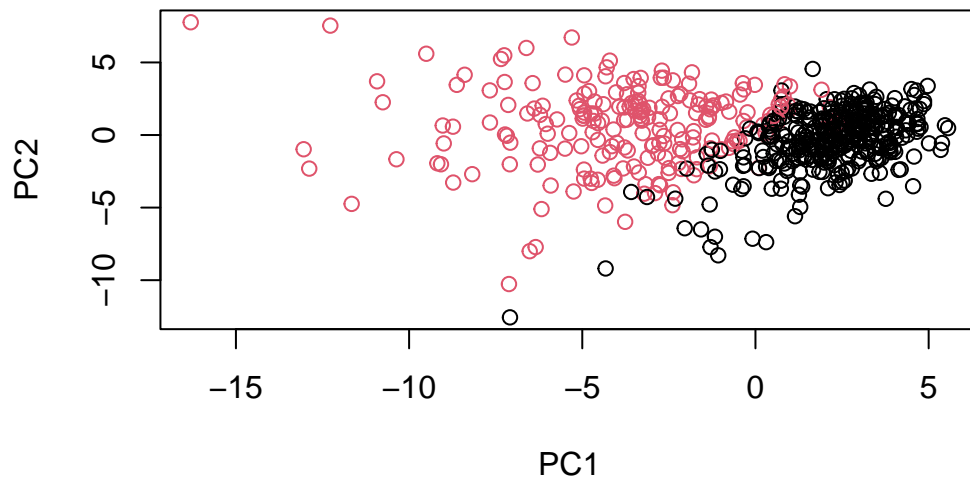


```
biplot(wisc.pr)
```



What stands out about the plot is that there are two main grouping representing the malignant and benevolent tumors in the different colors. This graph is very difficult to read. There is too much overlap and noise because it shows all of the different rows at once. We are not able to see which values are which.

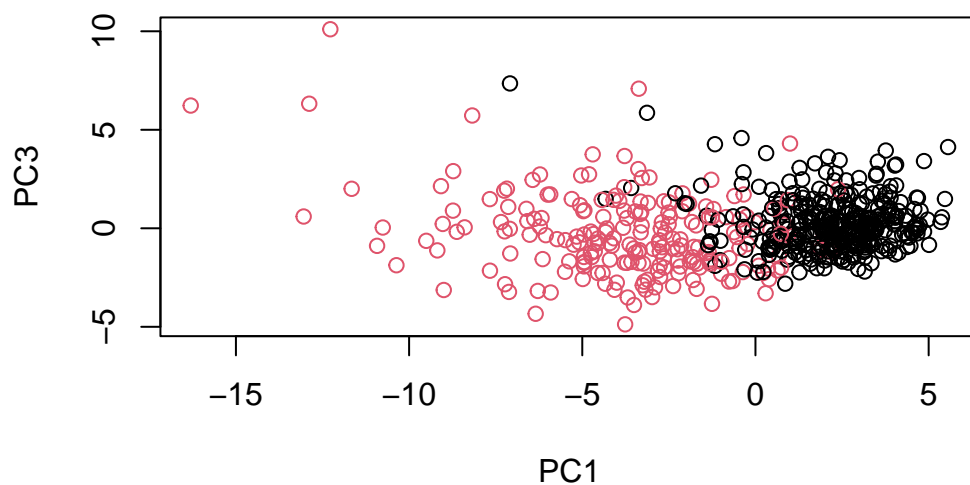
```
plot(wisc.pr$x[,1:2], col = diag, xlab = "PC1", ylab = "PC2")
```

Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

There are two main clusters representing the benign and malevolent tumors. It is a lot more organized and we can more easily see what is going on.

```
plot(wisc.pr$x[,1], wisc.pr$x[,3], col = diag, xlab = "PC1", ylab = "PC3")
```




Next, we will basically recreate this but use ggplot2.

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col= diagnosis) +
  geom_point()
```



Variance Explained

First we will calculate the variance explained by each principal component.


```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

Next, I am calculating the variance explained by each principal component.

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



Communicating PCA Results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC.

```
sorted <- sort(wisc.pr$rotation[,1])
#barplot(sorted)
sorted
```

concave.points_mean	concavity_mean	concave.points_worst
-0.26085376	-0.25840048	-0.25088597
compactness_mean	perimeter_worst	concavity_worst
-0.23928535	-0.23663968	-0.22876753
radius_worst	perimeter_mean	area_worst
-0.22799663	-0.22753729	-0.22487053
area_mean	radius_mean	perimeter_se
-0.22099499	-0.21890244	-0.21132592
compactness_worst	radius_se	area_se
-0.21009588	-0.20597878	-0.20286964
concave.points_se	compactness_se	concavity_se
-0.18341740	-0.17039345	-0.15358979

smoothness_mean	symmetry_mean	fractal_dimension_worst
-0.14258969	-0.13816696	-0.13178394
smoothness_worst	symmetry_worst	texture_worst
-0.12795256	-0.12290456	-0.10446933
texture_mean	fractal_dimension_se	fractal_dimension_mean
-0.10372458	-0.10256832	-0.06436335
symmetry_se	texture_se	smoothness_se
-0.04249842	-0.01742803	-0.01453145

-0.26085376

3. Hierarchical Clustering

Here, I am scaling the wisc.data data

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```

Now, we need to calculate the Euclidean distances between all of the pairs of observations in the data set we just scaled.

```
#dist() function shows all the Euclidean distances.
data.dist <- dist(data.scaled)
#data.dist
```

Now, we need to create a hierarchical clustering model using the complete linkage. We will apply the hclust() argument and assign this to wisc.hclust

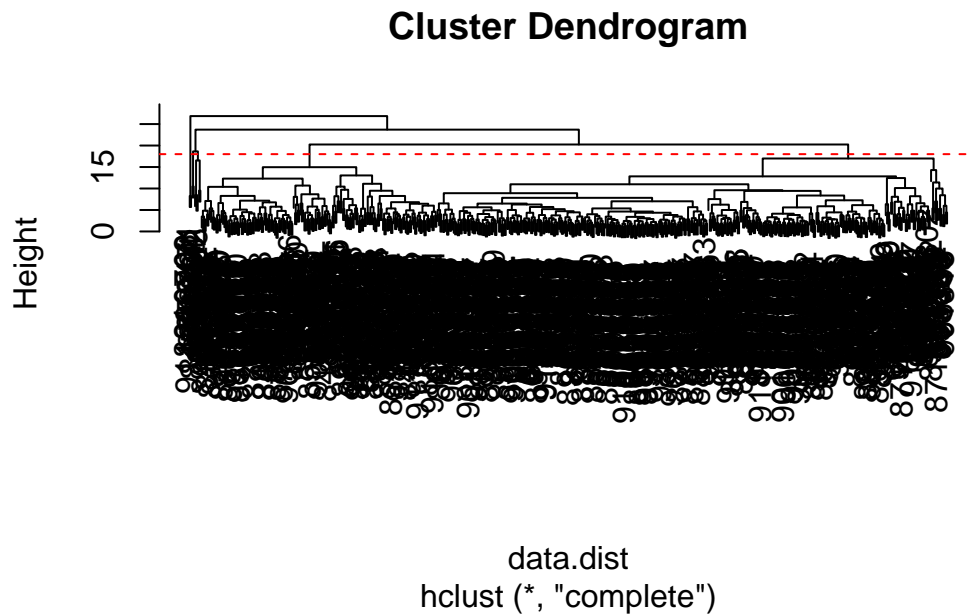
```
wisc.hclust <- hclust(data.dist, method = "complete")
#wisc.hclust
```

Results of Hierarchical Clustering

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

The height at which the clustering model has 4 groups is $h = 18$ (see code and graph)

```
plot(wisc.hclust)
abline(wisc.hclust, col = "red", lty = 2, h = 18)
```



Using Different Methods


Q12. Which method gives your favorite results for the same `data.dist` dataset? Explain your reasoning.

Let's try out the different results. We tried "complete" before, so now let's try "average" and "ward.D2"

This is what average looks like:

```
#average
wisc.hclust.average <- hclust(data.dist, method = "average")
plot(wisc.hclust.average)
```


Cluster Dendrogram



```
data.dist  
hclust (*, "average")
```

This is what ward.D2 looks like:

```
#ward.D2  
wisc.hclust.ward.D2 <- hclust(data.dist, method = "ward.D2")  
plot(wisc.hclust.ward.D2)
```



```
data.dist
hclust (*, "ward.D2")
```

I think I like ward.D2 the best as well. The other ones look less organized and you have to stare at them for longer in order to figure out what is going on since there are so many branches that go off from the top into other groups. ward.D2 on the other hand has one large, main branch at the top that separates into two obvious groups. It is a lot nicer for pattern recognition and feels more organized.


4. Combining Methods

Clustering on PCA Results

We will need to create a hierarchical clustering model using `method = "ward.D2"`.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method = "ward.D2")
plot(wisc.pr.hclust)
```


Cluster Dendrogram



```
dist(wisc.pr$x[, 1:7])  
hclust (*, "ward.D2")
```

Let's find out if these two groups of clusters in this dendrogram are malignant or benign:

```
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

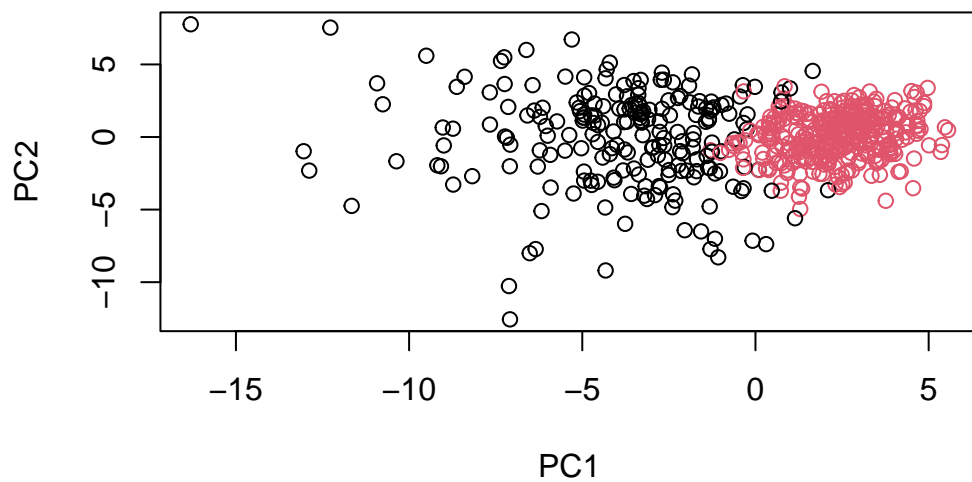
```
grps  
  1   2  
216 353
```

```
table(grps, diagnosis)
```


```
      diagnosis  
grps   B    M  
  1   28 188  
  2  329   24
```

To have a visual representation, let's make a plot where the two different groups are shown in different colors, black and red.

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diag)
```



Now, let's cut the hierarchical clustering model into 2 clusters and assign the results to `wisc.pr.hclust.clusters`

```
wisc.pr.hclust.clusters<- cutree(wisc.pr.hclust, k=2)

table(wisc.pr.hclust.clusters, diag)
```

```
      diag
wisc.pr.hclust.clusters  B  M
      1  28 188
      2 329  24
```

Q13. How well does the newly created model with four clusters separate out the two diagnoses?

I think it did a pretty good job for two clusters. In cluster 1, there are 28 benevolent diagnoses and 188 malignant, so it is mostly malevolent. In the second cluster, there are 329 benevolent diagnoses and 24 malignant diagnoses. There is a majority of one diagnoses in each and not too many points that are far off.

For four clusters, I will use the `table()` function to compare:

```
wisc.pr.hclust.clusters.4 <- cutree(wisc.hclust, k=4)

table(wisc.pr.hclust.clusters.4, diag)
```

```
      diag
wisc.pr.hclust.clusters.4  B  M
      1  12 165
      2   2   5
      3 343  40
      4   0   2
```

Again, it looks like it did a good job. Clusters 2 and 4 are very tiny though. I feel like the results were more accurate when `k=2` .

Q14. How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.pr.hclust.clusters.4, diag)
```

```

              diag
wisc.pr.hclust.clusters.4  B  M
              1  12 165
              2   2   5
              3 343  40
              4   0   2

```

...and compare that to the kmeans model for the clusters subset:

```

wisc.km.4 <- kmeans(wisc.data, centers = 4)
table(wisc.km.4$cluster, diag)

```

```

diag
  B  M
1  0 93
2  0 19
3 274  8
4  83 92

```

Before PCA, the hierarchical clustering model did not do as well in separating out the diagnoses. We definitely see better grouping when PCA is combined with hierarchical cluster modeling.

6. Prediction

```

#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

```

```

      PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
      PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
      PC15     PC16     PC17     PC18     PC19     PC20
[1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216

```

```


[2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
      PC21      PC22      PC23      PC24      PC25      PC26
[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27      PC28      PC29      PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820

```

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



- **Q16.** Which of these new patients should we prioritize for follow up based on your results?

Based on the results, we should prioritize patient 2 because this patient lies in the malevolent cluster and is therefore more likely to have a malevolent tumor that needs more rapid medical attention. Meanwhile, patient 1 lies in the benevolent cluster and is therefore less likely to need rapid medical attention if the clustering is accurate and the prediction holds true.