

# Class 12: Transcriptomics and the Analysis of RNA-Seq Data

Dahlia Loomis

5/12/23

First, we installed the DESeq package before moving on to the second portion of the lab.

## 2. Import countData and colData

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Now, we will have a brief look at the data

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG00000000003	723	486	904	445	1170
ENSG00000000005	0	0	0	0	0
ENSG00000000419	467	523	616	371	582
ENSG00000000457	347	258	364	237	318
ENSG00000000460	96	81	73	66	118
ENSG00000000938	0	0	1	0	2

	SRR1039517	SRR1039520	SRR1039521
ENSG00000000003	1097	806	604
ENSG00000000005	0	0	0
ENSG00000000419	781	417	509
ENSG00000000457	447	330	324
ENSG00000000460	94	102	74
ENSG00000000938	0	0	0

```
head (metadata)
```

```
      id      dex celltype    geo_id
1 SRR1039508 control  N61311 GSM1275862
2 SRR1039509 treated  N61311 GSM1275863
3 SRR1039512 control  N052611 GSM1275866
4 SRR1039513 treated  N052611 GSM1275867
5 SRR1039516 control  N080611 GSM1275870
6 SRR1039517 treated  N080611 GSM1275871
```

**Q1. How many genes are in this dataset?**

```
nrow(counts)
```

```
[1] 38694
```

There are 38694 genes in the dataset

**Q2. How many 'control' cell lines do we have?**

There are 4 control cell lines.

### 3. Toy Differential Gene Expression

First, we will get a better understanding of differential gene expression analysis.

```
control <- metadata[metadata[,"dex"]=="control",]
control.counts <- counts[ ,control$id]
control.mean <- rowSums( control.counts )/4
head(control.mean)
```

```
ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
      900.75           0.00           520.50           339.75           97.25
ENSG000000000938
      0.75
```

We will separate this code out to understand it better:

```
metadata[,"dex"] == "control"
```

```
[1] TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE
```

This table shows just the control samples

```
control <- metadata[metadata[, "dex"] == "control", ]
```

```
control
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
3	SRR1039512	control	N052611	GSM1275866
5	SRR1039516	control	N080611	GSM1275870
7	SRR1039520	control	N061011	GSM1275874

Now, we can separate out the id's

```
control$id
```

```
[1] "SRR1039508" "SRR1039512" "SRR1039516" "SRR1039520"
```

Now, we can use this as a subset in the counts data set and summarize the rows using the mean of the rows.

```
control.counts <- counts[, control$id]  
control.means <- rowMeans(control.counts)
```

**Q3. How would you make the above code in either approach more robust?**

We can use the rowMeans function. This allows the function to be applied to other datasets without having to alter the value that you are dividing by. In this particular data set, we divide by 4 since there are 4 genes, but this is not the case for all datasets. The number of values in each row is going to be different.

**Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug samples and assign to a labeled vector called treated.mean)**

```
treated <- metadata[metadata[, "dex"] == "treated", ]  
treated$id
```

```
[1] "SRR1039509" "SRR1039513" "SRR1039517" "SRR1039521"
```


```
treated.counts <- counts[,treated$id]
treated.means <- rowMeans(treated.counts)

#View(as.data.frame(treated.means))

meancounts <- data.frame(control.means, treated.means)
```

**Q5 (a).** Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look like the following.

```
plot(meancounts)
```




Now, we can make the plot a little fancier by using ggplot2.

**Q5(b).** You could also use the ggplot2 package to make this figure producing the plot below. What geom\_?() function would you use for this plot?

You would use the geom\_point function.

```
library(ggplot2)
ggplot(counts, aes(control.means, treated.means)) + geom_point(alpha = 0.5)
```




**Q6.** Try plotting both axes on a log scale. What is the argument to `plot()` that allows you to do this?

The argument we can use is `log()`

Note: this emits the 0s

```
plot(log(meancounts))
```



To calculate the log2 of the fold change between treated and control...

We need to add a log2 column to the data frame

```
meancounts$log2fc <- log2(meancounts[, "treated.means"]/meancounts[, "control.means"])
head(meancounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

To remove zero values...

```
zero.vals <- which(meancounts[, 1:2]==0, arr.ind=TRUE)
to.rm <- unique(zero.vals[, 1])
mycounts <- meancounts[-to.rm,]
```

This gives us the row and the column where we have a 0.

```
#which(meancounts[,1:2] == 0, arr.ind = TRUE)
```

**Q7. What is the purpose of the `arr.ind` argument in the `which()` function call above? Why would we then take the first column of the output and need to call `unique()` function?**

The `arr.ind` argument will tell us the row and column numbers for where we have 0s in the data frame instead of just getting a matrix of TRUE and FALSE. We then need to use the `unique()` function to avoid repetitions.

### Overexpressed and underexpressed genes

```
up.ind <- mycounts$log2fc > 2  
down.ind <- mycounts$log2fc < (-2)  
table(up.ind)
```

```
up.ind  
FALSE  TRUE  
21567   250
```

```
table(down.ind)
```

```
down.ind  
FALSE  TRUE  
21450   367
```

**Q8. Using the `up.ind` vector above, can you determine how many up regulated genes we have at the greater than 2 fc level?**

We have 250 genes.

**Q9. Using the `down.ind` vector above, can you determine how many down regulated genes we have at the greater than 2fc level?**

We have 367 genes.

**Q10. Do you trust these results? Why or why not?**

No. We would need to see the p-values first since we need some quantitative analysis to see if our results are significant. When we actually run the DESeq analysis, we find that the boundaries for the fold-change are much smaller.

## 4. DESeq2 Analysis

First, we need to load in the package

```
library(DESeq2)
```

Loading required package: S4Vectors

Loading required package: stats4

Loading required package: BiocGenerics

Attaching package: 'BiocGenerics'

The following objects are masked from 'package:stats':

IQR, mad, sd, var, xtabs

The following objects are masked from 'package:base':

anyDuplicated, aperm, append, as.data.frame, basename, cbind,  
colnames, dirname, do.call, duplicated, eval, evalq, Filter, Find,  
get, grep, grepl, intersect, is.unsorted, lapply, Map, mapply,  
match, mget, order, paste, pmax, pmax.int, pmin, pmin.int,  
Position, rank, rbind, Reduce, rownames, sapply, setdiff, sort,  
table, tapply, union, unique, unsplit, which.max, which.min

Attaching package: 'S4Vectors'

The following objects are masked from 'package:base':

expand.grid, I, unname

Loading required package: IRanges

Loading required package: GenomicRanges



Loading required package: GenomeInfoDb

Loading required package: SummarizedExperiment

Loading required package: MatrixGenerics

Loading required package: matrixStats

Attaching package: 'MatrixGenerics'

The following objects are masked from 'package:matrixStats':

colAlls, colAnyNAs, colAnys, colAvgsPerRowSet, colCollapse,  
colCounts, colCummaxs, colCummins, colCumprods, colCumsums,  
colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,  
colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,  
colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,  
colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,  
colWeightedMeans, colWeightedMedians, colWeightedSds,  
colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgsPerColSet,  
rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,  
rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,  
rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,  
rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,  
rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,  
rowWeightedMads, rowWeightedMeans, rowWeightedMedians,  
rowWeightedSds, rowWeightedVars

Loading required package: Biobase

Welcome to Bioconductor

Vignettes contain introductory material; view with  
'browseVignettes()'. To cite Bioconductor, see  
'citation("Biobase")', and for packages 'citation("pkgname")'.

Attaching package: 'Biobase'

The following object is masked from 'package:MatrixGenerics':

```
rowMedians
```

The following objects are masked from 'package:matrixStats':

```
anyMissing, rowMedians
```

```
citation("DESeq2")
```

To cite package 'DESeq2' in publications use:

Love, M.I., Huber, W., Anders, S. Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2 Genome Biology 15(12):550 (2014)

A BibTeX entry for LaTeX users is

```
@Article{,
  title = {Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2},
  author = {Michael I. Love and Wolfgang Huber and Simon Anders},
  year = {2014},
  journal = {Genome Biology},
  doi = {10.1186/s13059-014-0550-8},
  volume = {15},
  issue = {12},
  pages = {550},
}
```

Let's generate the specific object DESeq2 needs:

```
dds <- DESeqDataSetFromMatrix(countData = counts, colData = metadata, design = ~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

```
dds
```

```
class: DESeqDataSet
dim: 38694 8
metadata(1): version
assays(1): counts
rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
               ENSG00000283123
rowData names(0):
colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
colData names(4): id dex celltype geo_id
```

## DESeq analysis

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
results(dds)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 38694 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.1942	-0.3507030	0.168246	-2.084470	0.0371175
ENSG000000000005	0.0000	NA	NA	NA	NA

ENSG00000000419	520.1342	0.2061078	0.101059	2.039475	0.0414026
ENSG00000000457	322.6648	0.0245269	0.145145	0.168982	0.8658106
ENSG00000000460	87.6826	-0.1471420	0.257007	-0.572521	0.5669691
...	...	...	...	...	...
ENSG00000283115	0.000000	NA	NA	NA	NA
ENSG00000283116	0.000000	NA	NA	NA	NA
ENSG00000283119	0.000000	NA	NA	NA	NA
ENSG00000283120	0.974916	-0.668258	1.69456	-0.394354	0.693319
ENSG00000283123	0.000000	NA	NA	NA	NA
	padj				
	<numeric>				
ENSG00000000003	0.163035				
ENSG00000000005	NA				
ENSG00000000419	0.176032				
ENSG00000000457	0.961694				
ENSG00000000460	0.815849				
...	...				
ENSG00000283115	NA				
ENSG00000283116	NA				
ENSG00000283119	NA				
ENSG00000283120	NA				
ENSG00000283123	NA				

```
#always use the padj column for analysis
```

With the summary command, we can have a quick overview of the results

```
res <- results(dds)
summary(res, alpha = 0.05)
```

```
out of 25258 with nonzero total read count
adjusted p-value < 0.05
LFC > 0 (up)      : 1242, 4.9%
LFC < 0 (down)    : 939, 3.7%
outliers [1]      : 142, 0.56%
low counts [2]    : 9971, 39%
(mean count < 10)

[1] see 'cooksCutoff' argument of ?results
[2] see 'independentFiltering' argument of ?results
```

## 6. Data Visualization

### Volcano Plots

```
plot(res$log2FoldChange, -log10(res$padj))
```

